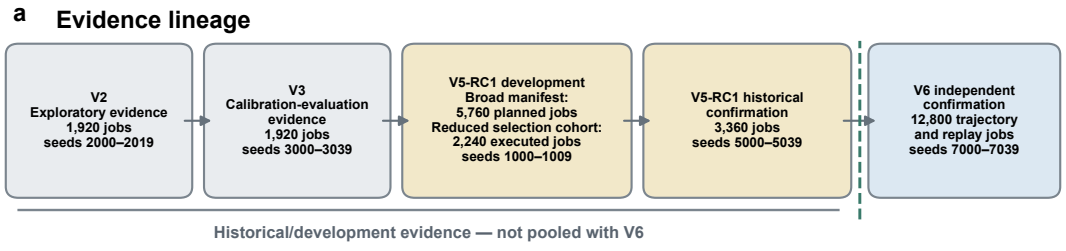
a
Evidence lineage
V2
Exploratory evidence
1,920 jobs
seeds 2000–2019
V3
Calibration-evaluation evidence
1,920 jobs
seeds 3000–3039
V5-RC1 development
Broad manifest:
5,760 planned jobs
Reduced selection cohort:
2,240 executed jobs
seeds 1000–1009
V5-RC1 historical confirmation
3,360 jobs
seeds 5000–5039
V6 independent confirmation
12,800 trajectory and replay jobs
seeds 7000–7039
Historical/development evidence — not pooled with V6
b
Prospectively locked V6 confirmation design
4 Matbench tasks
2 split types
2 tree ensembles
16 task × split × model environments
40 paired seed blocks
(seeds 7000–7039)
c
Separate registered outcomes and families
L1: campaign MACE
lower is better
H1, H5A, H6A, H6B
L4: top-1% recall nAUC
higher is better
H2A, H2B, H5B, H5C, H3, H4
d
Locked inference and interpretation boundary
Within each seed:
equal mean across
16 environments
10,000 paired bootstrap
10,000 sign flips
Family-specific
BH correction
Environment effects:
heterogeneity only